

Advanced Computer Graphics

Oliver Eisenberg

December 11, 2019

1 Labs

1.1 Lab 1 - Rasterising Lines

The line drawing function was updated to use an integer based approach. Originally the algorithm was used for positive x,y values and then altered by mirroring the result onto the other three quadrants. This created the following image *ref*.

This seems correct, however, on further inspection it was realised that the lines seemed to have a double thickness. This was thought to have occurred as the conditions to be in what quadrant was changing as the line was drawn. Therefore the process was reattempted in an effort to improve and refactor code.

The result was the correct output as shown here. This was done by .. Code was refactored into a single x and y direction methods where an initial function call was responsible for determining when to call the relevant one. By calculating directions based off of the start and end coordinates.

1.2 Lab 2 - Reading Models

The given code in polymesh.cpp was updated to be able to read 3D objects and generate wireframes.

1.3 Lab 3 - Simple Raytracing

1.3.1 Raycasting

1.3.2 Triangle intersection

To compute triangle intersections the Möller–Trumbore algorithm was used. This was used instead of the method on the *slides* anticipating the requirements for the barycentric coordinates

to complete Gouraud shading further in the coursework.

1.4 Lab 4 - Basic Lighting and Shadows

Starting to work on lighting prompted further refactoring. In doing so, a scene class was created to create and store all the objects and lighting for the render. All lights belong to the base light class, this allows to predetermine shared variables and functions a light should have.

1.4.1 Spotlights

1.4.2 Pointlights

The slides refer to two methods to create pointlights, with and without an associated direction. As spot lights are directional, pointlights with a constant intensity were implemented. This allows a pointlight to be placed between objects to cast shadow outwards.

1.4.3 Shadows

Shadows are checked by sending a ray from the hit position to each light in the scene. If an intersection occurs the intensity from that light is removed. This allows for shadows of different intensity depending on the source.

2 Optimisations

Coloured Diffuse Originally objects had three intensity values representing the three primary colours in addition to the diffuse and specular coefficients. This was later updated to use a material and colour class.

By splitting intensity into colours it allows the addition of coloured lights

and control over what colour objects can reflect.

3 Advanced Features

3.1 Photon Mapping

Some difficulty was encountered when choosing static libraries, this heavily influenced the chosen library to handle KD Trees. The library chosen is *Alglib* as it was written to be added like normal classes where you include the relevant headers and compile the cpp files.

3.1.1 Random emission - Lighting

Lights had to be updated with relevant random emission direction and position functions. Depending on the light

3.1.2 Direct lighting




Figure 1: Render with direct diffuse lighting only

3.1.3 Specular & Gloss

This was straight forward to implement as, due to the intensity, it is rendered using the existing raytracing method that was developed during the lab handins.

3.1.4 Caustics

Caustics are computationally intensive to calculate using standard raytracing methods but was straight forward to implement using a separate caustic photon map. In a caustic map photons are emitted from the light source to reflective and transmissive objects, later when rendering the caustic map is calculated using a radiance estimate. A caustic map can be filtered prior to rendering to smooth out hard edges - this hasn't been done in my implementation due to time constraints.




Figure 2: Render with caustics